

# Concurrency

---

Lukas Prokop

August 5, 2020

RustGraz community



# Prologue

---



# Rusty Days

Cause some designs are just ahead of time...



- Webference organized by **Rust Wrocław**
- 8 talks between 2020-07-27 and 2020-07-31
- I will give a short talk summary

**Mon, 2020-07-27 — Steve Klabnik**  
**“Should we have a Rust 2021 edition?”**



## Recap: Rust 2021 edition Talk

rustc is a multi-pass compiler: AST  $\rightarrow$  HIR  $\rightarrow$  MIR  $\rightarrow$  LLVM-IR

**HIR** expand source code to simpler primitive statements  
(type checking, method lookup)

**MIR** is all about control flow (control-flow, borrow checking)

- *lowering*: “The MIR is then lowered to LLVM-IR”
- “cargo check” omits code generation step
- *goal*: “query-based” compiler (memoized answers)
- editions are not allowed to differ in MIR

**nightly release** every night

**stable + beta release** every six weeks



### Editions:

- **is** longer-term progress, breaking changes, new editions are opt-in, can't change everything
- **can** can introduce new keywords, repurpose syntax (e.g. deprecate trait, introduce dyn trait)
- **cannot** cannot change coherence rules, standard library
- when do editions happen? no policy
- edition 2018 (= rust 1.31) was “a really big project for the various teams” and “a lot of burnout amongst contributors” (“I was a total mess”)
- editions feature-driven or time-boxed? time-boxed! 3 years is a nice compromise between yearly and 5-year releases
- Steve: We should have a 2021 edition. Much smaller than 2018.



## Recap: Rust 2021 edition Talk

```
meisterluk@gardner ~ % cat test.rs
fn async() -> u64 {
    42
}
```

```
fn main() {
    println!("{}", async());
}
```

```
meisterluk@gardner ~ % rustc test.rs
```



## Recap: Rusty days

```
meisterluk@gardner ~ % rustc --edition=2018 test.rs
error: expected identifier, found keyword `async`
--> test.rs:1:4
  |
1 | fn async() -> u64 {
  |      ^^^^^ expected identifier, found keyword
  |
help: you can escape reserved keywords
      to use them as identifiers
  |
1 | fn r#async() -> u64 {
  |      ^^^^^^^
```



**Mon, 2020-07-27 — Michalina Kotwica**  
**“Low-level optimization of algebraic and similar  
structures”**



## Recap: Low-level optimization Talk

- type algebra
- unit & never type
- memory layout (ABI) of composite types (struct, enum, futures)
- struct Bar(u16, u8, u16, u8):  
C: 11 11 33 xx 22 22 44 xx Rust: 11 11 22 22  
33 44 (fields reordered)
- **Option** uses a discriminator, **Option**<**Option**<**u64**>> uses only one discriminator
- 41.2 % of enums have no type arguments, 16.4 % of enums have one type argument, 23.7 % of enums have two type arguments

**Tue, 2020-07-28 — Peter Parkanyi**  
**“Fast encrypted backups with Rust - ‘How I  
stopped worrying and love mmap’”**



## Recap: Fast encrypted backups Talk

<https://github.com/rsdy/zerostash>

- End-to-End Encryption
- latency and throughput
- Zero-metadata data stash: deduplicated, works as a file system and key/value store
- Cryptographic primitives:
  - Passwords** Argon2
  - Indexing** Blake2
  - Compression** LZ4
  - Encryption** ChaCha20-Poly1305
  - Deduplication** SeaHash
- Profiling: perf on Linux, Instruments on macOS
- mmap versus read

**Wed, 2020-07-29 — Lachezar Lechev**

**“Drone Control - ‘Controlling a drone using Rust over WiFi’”**



## Recap: Drone Control Talk

<https://github.com/AeroRust/Welcome>

- TCP connections Handshake, establish connection
- JSON requests & responses
- ping-pong within 7s
- **scroll** crate: “A suite of powerful, extensible, generic, endian-aware Read/Write traits for byte buffers”

**Wed, 2020-07-29 — Nell Shamrell - Harrington**  
**“The Rust Borrow Checker - A Deep Dive”**



## Recap: Drone Control Talk

### Slides on slideshare

- “Is the Borrow Checker a friend or a foe?”
- Stages of Compilation:
  1. Lexical Analysis
  2. Parsing
  3. Semantic Analysis (Borrow Checker!)
  4. Optimization
  5. Code Generation
- BC tracks initializations/moves and applies lifetime inference
- Lifetime of a variable has two definitions
  - “Span of time before the value of a variable gets freed”
  - “scope of a variable”
- “If you make a reference to a value, the lifetime of that reference cannot outlive the scope of the value”
- <https://rustc-dev-guide.rust-lang.org/>



**Thu, 2020-07-30 - Jan-Erik Rediger**  
**“Leveraging Rust to build cross-platform mobile  
libraries”**



## Recap: Mobile Libraries Talk

### Slides on slideshare

- [Firefox Telemetry project](#)
- Collect performance metrics for our products, package pings at controlled schedules
- Three Principles
  - Stay Lean
  - Build Security
  - Engage Your Users
- Telemetry *scalars*: `Scalars.yaml` (metadata: `bug_numbers`, `description`, `expires`, `notification_emails`, ...)



## Recap: Mobile Libraries Talk

- Glean core → Glean FFI → Glean Kotlin/Swift → Android/iOS app
- **cbindgen** crate: “A tool for generating C bindings to Rust code.”
- 10 Glean compilation targets supported
- Hello World with JNI (also see Otavio Pace’s talk)
- tagged unions are generated by bindgen for rust enums
- ProtoBuf to serialize data
- R8 minifies/optimizes JVM bytecode (successor to proguard)
- here: only invoke Rust code from Kotlin (never the other way around)

**Fri, 2020-07-31 - Luca Palmieri**

**“Are we observable yet? Telemetry for Rust APIs  
- metrics, logging, distributed tracing”**



## Recap: Telemetry for Rust APIs Talk

- Developer of **DonateDirect**
- New project: fast versus reliable (metrics, tracing, logs, ...)
- Claim: convenience beats correctness
- Metrics give us an aggregate picture of the system state
- **actix\_web\_prom** crate provides a pluggable middleware with standard Prometheus metrics out of the box
- **log** crate and **tracing** crate to dump structured logging with JSON output and forward spans to **elasticsearch** and then **kibana**



### Key takeaways:

- Lack of telemetry is a ticking bomb
- Diagnostic instrumentation has to be easy
- Metrics to alert and monitor system state
- High cardinality is key to being able to detect and triage unknown unknowns
- Span as unit of work abstraction
- You must be able to trace a request across different services

**Fri, 2020-07-31 — Tim McNamara**

**“How 10 open source projects manage unsafe code”**



## Recap: Managing unsafe code Talk

### Unsafe guidelines for the impatient:

- Use `#[deny(unsafe_code)]`
- Add comments to all unsafe blocks
- Let someone read the unsafe block comment.  
If they cannot explain afterwards, revise the comment

### Remarks:

- Question unsafe! Look for safe alternatives
- rust std: “Unsafe code block need a comment explaining why they’re ok”
- “We try to create a situation where we, as a team, are building safe software and we are mentally switched on if we go to the unsafe module”
- [UCG WG - Rust's Unsafe Code Guidelines Working Group](#)





## Recap: Managing unsafe code Talk

- `actix_web` `incident`
- `#[deny(unsafe_code)]` and `#[allow(unsafe_code)]` (can be nested)
- `cargo-geiger` “detects usage of unsafe Rust in a Rust crate and its deps”, introduces `#![forbid(unsafe_code)]`
- `exa` uses syscalls natively, `BLAKE3` uses SIMD instructions, `Firecracker` interacts with a hypervisor, `librsvg` talks to GLib, `toolshed` deals with pointers in one module, `terminusdb` interacts with Prolog, `Fuchsia OS` interacts with the non-rust kernel

# Dialogue

---

# Concurrency



## Definition: concurrency

*In computer science, concurrency is the ability of different parts or units of a program, algorithm, or problem to be executed out-of-order or in partial order, without affecting the final outcome.*

—*Concurrency (computer science)*



## Definition: parallelism

Parallel computing *is a type of computation in which many calculations or the execution of processes are carried out simultaneously.*

—*Parallel computing*



Formally, different models exist:

- Parallel random-access machine
- Actor model
- Petri nets
- Process calculi (CCS, CSP,  $\pi$ -calculus)
- Tuple spaces
- Simple Concurrent Object-Oriented Programming (CSOOP)
- Reo Coordination Language

... to design distributed systems.



## Models of concurrency

For example, Erlang is famous for its distributed computational model.

But, in rust, we stick to a von-Neumann-like model:

- Instructions are executed in order.
- We have stacks, heaps, bss and data sections to organize the memory.
- We make syscalls to the kernel and compile against some ISA.

Thus it is a question of an API. We define *concurrent units* and they might run in *parallel*.



## Definition: concurrency

Concurrency as a question of level of granularity:

1. Instruction
2. blocks of instructions
3. Function
4. Several stacks, one heap
5. Process

In general: performance as incentive.





## Problems of concurrency

### 1. Data dependency (synchronization problem)

**let** A and B be two concurrent units.

Both want to increment  $x$

**A** reads that  $x$  is 41

**B** reads that  $x$  is 41

**A** increments  $x$  and gets  $y := 42$

**B** increments  $x$  and gets  $z := 42$

**A** writes  $x := y$

**B** writes  $x := z$

**x** is 42

### 2. Concurrent units need to exchange messages

### 3. Concurrent unit waits for an event

### 4. Interrupt concurrent unit (*preemption*)

### 5. Determine concurrent unit has finished

**In rust (from low-level to high-level)**

**On instruction-level**



# SIMD instructions

SIMD = **s**ingle **i**nstruction, **m**ultiple **d**ata

Run one instruction, apply arithmetic/logic/data-handling/memory instructions to several values simultaneously. 8/16/32/64/128/512

**RISC-V** I prefer to talk about the RISC-V ISA, but RISC-V basically dropped its **Packed SIMD** extension and develops a new “**P**” **extension**. Thus, I switch to x86\_64.

**x86\_64** **Streaming SIMD Extensions (SSE)** and **Advanced Vector Extensions (AVX)**



Operands (put into XMM/YMM registers):

**SSE** four f32

**SSE2** two f64, two i64, four i32, eight i16, sixteen u8

**SSE3** (only 13 new instructions)

**SSE4** (only 54 new instructions)

**AVX** eight f32, four f64

**AVX2** 256-bit registers for almost everything

**AVX-512** 512-bit registers but instructions split into multiple extensions



## SIMD support

```
% cat /proc/cpuinfo | grep "sse\|avx"
flags      : fpu ... sse sse2 ... ssse3 ... sse4_1 sse4_2
            ... avx ... avx2 ... flush_l1d
flags      : fpu ... sse sse2 ... ssse3 ... sse4_1 sse4_2
            ... avx ... avx2 ... flush_l1d
flags      : fpu ... sse sse2 ... ssse3 ... sse4_1 sse4_2
            ... avx ... avx2 ... flush_l1d
flags      : fpu ... sse sse2 ... ssse3 ... sse4_1 sse4_2
            ... avx ... avx2 ... flush_l1d
```



## (No) SIMD instructions

```
fn foo(a: &[u8], b: &[u8], c: &mut [u8]) {  
    for ((a, b), c) in a.iter().zip(b).zip(c) {  
        println!("{}", {} {}, a, b, c);  
        *c = *a + *b;  
    }  
}
```

```
fn main() {  
    let a: [u8; 4] = [0xDE, 0xAD, 0xBE, 0xEF];  
    let b: [u8; 4] = [0x00, 0x01, 0x02, 0x03];  
    let mut c = [0u8; 4];  
    foo(&a, &b, &mut c);  
  
    println!("{}", c);  
}
```



## (No) SIMD instructions

On **godbolt** with `-C opt-level=1` (or 0):

```
.LBB58_2:
```

```
; ...
```

```
mov     rdx, qword ptr [rsp + 24]
```

```
movzx   ecx, byte ptr [rcx]
```

```
add     cl, byte ptr [rax]
```

```
mov     byte ptr [rdx], cl
```

```
mov     rdi, r14
```

```
; ...
```

```
jne     .LBB58_2
```





## SIMD instructions

On **godbolt** with `-C opt-level=2` (or higher):

```
movdqu    xmm2, xmmword ptr [rdx + rcx + 32]
paddb     xmm2, xmm0
movdqu    xmm0, xmmword ptr [rdx + rcx + 48]
paddb     xmm0, xmm1
movdqu    xmmword ptr [r8 + rcx + 32], xmm2
```

**paddb**: “add packed integer” instruction



## SIMD instructions

via `std::arch` (“SIMD and vendor intrinsics module”):

```
#[cfg(all(
    any(target_arch = "x86", target_arch = "x86_64"),
    target_feature = "avx2"
))]
fn foo() {
    #[cfg(target_arch = "x86")]
    use std::arch::x86::_mm256_add_epi64;
    #[cfg(target_arch = "x86_64")]
    use std::arch::x86_64::_mm256_add_epi64;

    unsafe {
        _mm256_add_epi64(...);
    }
}
```



## SIMD summary

- rust uses the LLVM stack, which implements *auto-vectorization*
- usually, we get SIMD instructions for free
- sometimes, you want to use the explicitly
- you can always use inline assembly with `unsafe`  
`(x86::time::rdtsc)`
- `packed_simd` provides a high-level API
- also: `faster`, `ssimd`
- mostly, libraries are specific for an application domain, e.g.:
  - `numeric-array` crate
  - `directx_math` crate
  - `pqcrypto-classicmceliece` crate

**On blocks of code**

**On OS thread level**



# threads

Copy stack (expensive?), share heap.



- “A mutual exclusion primitive useful for protecting shared data”
- A mutex can be `locked` (true/blocked) or released
- Also `try_lock` (true/false)
- Only one concurrent unit can hold a lock simultaneously
- If the mutex lock goes out of scope, another unit can acquire the lock.



## Mutex

```
use std::sync::Mutex;

fn main() {
    let val = Mutex::new(0u8);
    {
        *val.lock().unwrap() += 1;
        *val.lock().unwrap() += 1;
        *val.lock().unwrap() += 1;
    }
    println!("{}", val.lock().unwrap());
}
```





## Mutex

```
use std::sync::{Arc, Mutex};
use std::thread;

fn main() {
    let mutex = Arc::new(Mutex::new(0));
    let c_mutex = mutex.clone();

    thread::spawn(move || {
        *c_mutex.lock().unwrap() = 10;
    }).join().expect("thread::spawn failed");
    println!("{}", *mutex.lock().unwrap());
}
```

**async & await**



## async & await

Similar to C# and JavaScript. `async` keyword.

```
async fn sub() -> u8 {  
    42u8  
}
```

```
fn main() {  
    println!("{}", sub())  
}
```



## async & await

```
meisterluk@gardner ~ % rustc --edition=2018 async1.rs
error[E0277]: `impl std::future::Future`
             doesn't implement `std::fmt::Display`
--> async1.rs:6:20
   |
6 |     println!("{}", sub())
   |                    ^^^^^^ `impl std::future::Future`
   |                    cannot be formatted with the default formatter
   |
= help: the trait `std::fmt::Display` is not implemented for
`impl std::future::Future`
= note: in format strings you may be able to use `{:?}`
(or `{:#?}` for pretty-print) instead
= note: required by `std::fmt::Display::fmt`
= note: this error originates in a macro (in Nightly builds,
run with -Z macro-backtrace for more info)
```

error: aborting due to previous error

For more information about this error, try `rustc --explain E0277`.



*A **future** represents a value, that is not ready yet. Eventually, the future resolves to a value.*

*—**withoutboats** in a Rust Latam talk 2019*



## async & await

await keyword. Not await X, but X.await.

```
async fn sub() -> u8 {  
    42u8  
}
```

```
fn main() {  
    println!("{}", sub().await)  
}
```



# async & await

```
meisterluk@gardner ~ % rustc --edition=2018 async2.rs
error[E0728]: `await` is only allowed inside `async`
            functions and blocks
```

```
--> async2.rs:6:20
   |
5 | fn main() {
   |     ---- this is not `async`
6 |     println!("{}", sub().await)
   |                        ^^^^^^^^^^^^^^^^^ only allowed
   |                        inside `async` functions and blocks
```

```
error: aborting due to previous error
```

For more information about this error, try ``rustc --explain E0728``.



So who can call the first async function? the *executor* (scheduler)

- e.g. `async-std`, `smol`, `tokio`
- executor is thus exchangeable
- executor allocates memory per future
- futures are like state machines between states

the *reactor*:

- Executor: is future X ready to go? Yes, then go. Else:
- Reactor: I will take care of the future and report back when its ready





## async & await

```
use smol::io;
```

```
async fn sub() -> u8 {  
    42u8  
}
```

```
fn main() -> io::Result<()> {  
    smol::run(async {  
        println!("{}", sub().await);  
        Ok(())  
    })  
    // prints 42  
}
```



## async & await

```
async fn sub() -> u8 {  
    42u8  
}
```

```
fn main() {  
    async {  
        println!("{}", sub().await)  
    };  
}
```



## async & await

```
meisterluk@gardner ~ % rustc --edition=2018 async3.rs
warning: unused implementer of `std::future::Future`
        that must be used
```

```
--> async3.rs:6:5
   |
 6 | /      async {
 7 | |          println!("{}", sub().await)
 8 | |      };
   | |_____^
   |
= note: `#[warn(unused_must_use)]` on by default
= note: futures do nothing unless you `.await`
        or poll them
```

```
warning: 1 warning emitted
```

```
meisterluk@gardner ~ % ./async3      # prints nothing
```



Other async/await implementations:

- Calling an async function, schedules it
- Javascript calls *Promise*, what rust calls *Future*
- Green threads: a scheduler is compiled in every executable to manage small threads (M:N threading) (not OS threads!) (e.g. erlang, golang goroutines). rust 1.0 dropped green threads.

Rust implementations:

- Zero-cost: only await schedules the function (“lazy”)

**futures**, an official Rust crate that lives in the rust-lang repository A runtime of your choosing, such as Tokio, *async<sub>s</sub>td*, *smol*, etc.



# async & await

```
fn get_two_sites() {  
    // Spawn two threads to do work.  
    let thread_one = thread::spawn(|| download("https://foo.com"));  
    let thread_two = thread::spawn(|| download("https://bar.com"));  
  
    // Wait for both threads to complete.  
    thread_one.join().expect("thread one panicked");  
    thread_two.join().expect("thread two panicked");  
}
```

threading approach via [async-book](#)



# async & await

```
async fn get_two_sites_async() {  
    // Create two different "futures" which, when run to completion,  
    // will asynchronously download the webpages.  
    let future_one = download_async("https://foo.com");  
    let future_two = download_async("https://bar.com");  
  
    // Run both futures to completion at the same time.  
    join!(future_one, future_two);  
}
```

async approach via [async-book](#)



## async & await summary

- Async-await is a way to write functions that can "pause", return control to the runtime, and then pick up from where they left off
- proposed 2016, didn't make edition 2018, landed in rust 1.39
- Memory management
- [wasm-bindgen-futures](#) binds Rust Future to Javascript Promise

**On process-level**



**But how many units?**



# How many concurrent units?

**How many?** Difficult.

- CPU cycles per instruction differs, caches everywhere!
- Thus, **benchmark, benchmark, benchmark!**

But how?

- No stable stdlib support in rust!
- RFC 29553: Tracking issue for [bench] and benchmarking support
- e.g. **criterion** crate



# How many concurrent units?

## Heuristics:

- RFC 985: Add `std::env::concurrency_hint`
- `num_cpus` crate: `num_cpus::get()` ;
- `dupfiles-go` and `ripgrep` experience: the number of concurrent units reading files should roughly correspond to the number of logical CPUs

# Epilogue

---



## Quiz

**What does SIMD stand for?**

**What is a mutex?**

**What is a future?**

**Heuristically, what might be a good number of concurrent units?**



## Quiz

**What does SIMD stand for?**

Single Instruction, Multiple Data

**What is a mutex?**

**What is a future?**

**Heuristically, what might be a good number of concurrent units?**



## Quiz

**What does SIMD stand for?**

Single Instruction, Multiple Data

**What is a mutex?**

a data structure to give mutually exclusive access to a code section

**What is a future?**

**Heuristically, what might be a good number of concurrent units?**



## Quiz

**What does SIMD stand for?**

Single Instruction, Multiple Data

**What is a mutex?**

a data structure to give mutually exclusive access to a code section

**What is a future?**

represents a value, that is not ready yet

**Heuristically, what might be a good number of concurrent units?**





### **What does SIMD stand for?**

Single Instruction, Multiple Data

### **What is a mutex?**

a data structure to give mutually exclusive access to a code section

### **What is a future?**

represents a value, that is not ready yet

### **Heuristically, what might be a good number of concurrent units?**

number of logical CPUs of the machine



## Next time

Next meetup    Wed, 2020/08/26

Topic            I/O (files, file formats, simple TCP server)

**Thank you!**

